

On August 2, 2026, the EU AI Act's high-risk rules become binding.

KEY POINTS

- Penalty for prohibited practices: 35 million EUR or 7 percent global revenue, whichever is higher.
- Other high-risk violations: 15 million EUR or 3 percent.
- 50 percent of enterprises still have no AI inventory.
- 75 days remain. Article 11 technical files are not optional.
- Being caught flat-footed is a budget event, not a fine line item.

SPEAKER NOTES

Today is May 19, 2026. In 75 days, on August 2, the EU AI Act's high-risk system rules become legally binding across the European Union. Article 99 sets the ceiling at 35 million euros or 7 percent of global revenue for prohibited practices, and 15 million or 3 percent for other high-risk violations. The European Commission's own DG-CNECT survey says more than half of enterprises do not have a current inventory of the AI systems they deploy. They cannot classify what they cannot see. And they cannot produce an Article 11 technical file for a system they have not catalogued. That is the gap ComplyForge closes.

VISUAL DIRECTION

Full-bleed countdown bar in risk-prohibited red. Centered text: 75 DAYS on the first line, August 2, 2026 on the second, in accent amber. Sub-line in text-dim: "EU AI Act Article 6 + Annex III obligations binding". ---

Compliance teams have no tooling for Article 11.

KEY POINTS

- The buyers: CISO, General Counsel, Head of Compliance, AI Ethics Lead.
- The need: AI inventory, EU AI Act classification, Article 11 technical file, Article 27 FRIA, enforced policy, audit trail.
- Trapdoor Tech ships access control, not regulatory documents.
- Sky/AgentWatch ships detection observability, not an Annex IV draft.
- Harvey, Ironclad, PreMortem, Infinite Context — none cover Article 11 plus FRIA plus enforcement.

SPEAKER NOTES

There are good tools in adjacent lanes. Trapdoor Tech is contractor access control on top of Lobster Trap. Sky and AgentWatch are detection observatories. Harvey and Ironclad live in legal contracts. PreMortem AI does retrospective failure analysis. Infinite Context does multi-agent research. None of them produces an Annex IV technical file that an EU auditor would recognise. None of them issues a Lobster Trap policy that enforces what the file claims. None of them keeps an audit log a regulator can replay. ComplyForge is the only one that closes that loop end to end.

VISUAL DIRECTION

Six-row by four-column matrix table. Rows: AI inventory, EU AI Act classification, Article 11 PDF, Article 27 FRIA, Lobster Trap enforcement, Audit trail. Columns: Trapdoor Tech, AgentWatch, Harvey, ComplyForge. Each cell is a tick or empty circle. Only the ComplyForge column is full. ---

An autonomous agent that reasons, plans, and executes — in 60 seconds.

KEY POINTS

- Four artefacts per run: tier classification, Article 11 PDF plus FRIA, Lobster Trap YAML, audit log.
- Single async pipeline; classifier first, then DocAgent and PolicyAgent in parallel; PDF rendered at the end.
- Gemini 2.0 Flash for classification and structured documentation.
- Veea Lobster Trap for live policy enforcement on the wire.
- Locked at 1 plus 3 agents — not a 5-to-7 mesh.

SPEAKER NOTES

The architecture is deliberately small. One orchestrator. Three sub-agents. ClassifierAgent reads the agent descriptor and the full EU AI Act taxonomy and emits a tier with cited articles. DocAgent fans out ten concurrent Gemini calls, one per Article 11 section plus the FRIA, and assembles a multi-page PDF. PolicyAgent emits Veea Lobster Trap YAML in the real schema. The orchestrator runs DocAgent and PolicyAgent in parallel; the end-to-end pipeline finishes in a few seconds. We chose this shape on purpose: large agent meshes are in the trough of disillusionment because their error rates compound.

VISUAL DIRECTION

Block diagram. Top: AgentDescriptor input. Middle: a single ComplianceOrchestrator box. Three branches below: ClassifierAgent, DocAgent (annotated "9 Article 11 sections + 1 FRIA, async.gather"), and PolicyAgent (annotated "Veea Lobster Trap YAML"). Bottom: four output tiles — Risk tier, Article 11 PDF, Policy YAML, Audit log. Sponsor logos for Google Gemini and Veea Lobster Trap on the side. ---

Solo build, sponsor-native, real engineering honesty.

KEY POINTS

- 94 passing tests (87 unit plus 7 new endpoint tests); 1 opt-in end-to-end test.
- Verified Veea Lobster Trap schema differs from BUILD_BIBLE placeholder; patched in two hours and documented every field.
- Lobster Trap binary has no hot reload; we surface `manual_restart_required` in the API response rather than fake one.
- Two sponsor stacks, used as the manuals describe: Gemini for inference, Veea Lobster Trap MIT for enforcement.
- 1 plus 3 agents, not 5 to 7. Smaller surface, fewer compounding errors.

SPEAKER NOTES

Two things matter here. First, every claim on this slide is on disk. We have 94 passing tests today: 87 unit tests plus 7 new endpoint tests for the deploy-policy and inventory-zip endpoints, with one end-to-end test that is opt-in via the e2e marker. Second, we have not faked anything we could not verify. The Veea Lobster Trap policy schema in our brief turned out to differ from the real binary's loader. We read `serve.go` and the loader, patched our `PolicyAgent` in one centralised helper, and shipped a diff document. The same binary has no reload mechanism. We tell the operator that, with the policy file path, instead of inventing a flag.

VISUAL DIRECTION

Two-column layout. Left column: a small stat card grid — "94 tests passing", "5 enterprise agents seeded", "8 attack payloads", "Schema diff patched in 2 hours", each in bg-panel with the accent color number. Right column: a soft border-soft callout with the line "We will not fake what we cannot verify." in text-main, centred. ---

Watch high-risk to Article 11 to enforced to audited. Sixty seconds.

KEY POINTS

- Pick ResumeRanker from the seeded inventory.
- Classifier returns HIGH_RISK with Annex III(4) employment cited.
- DocAgent renders the Article 11 PDF and FRIA; PolicyAgent emits the 5-rule YAML.
- Deploy writes the policy to disk; restart Lobster Trap with the new file.
- Fire prompt-injection payload #2; proxy denies; audit log lights up.
- Ask: judges' vote, Veea Discord support, lablab community share.

SPEAKER NOTES

The demo is one click. Pick ResumeRanker, the HR resume screener. Watch four streaming steps: classify, generate documents, generate policy, render PDF. The classifier returns HIGH_RISK with Annex III(4). The DocAgent renders an Article 11 PDF with the nine Annex IV sections and the six-section FRIA. The PolicyAgent emits a five-rule Veea Lobster Trap policy that cites Articles 10, 12, 14, 15, and 72. Click Deploy. The API honestly tells us the binary needs a restart with the new policy file. We restart. We fire the prompt-injection attack. The proxy denies it under the injection rule. The dashboard's enforcement log lights up within two seconds. That is the loop. Vote ComplyForge.

VISUAL DIRECTION

Four-frame horizontal storyboard: (1) sample-agent grid with ResumeRanker selected and a ring-2 ring-accent; (2) the streaming StepProgress component mid-pipeline; (3) the Lobster Trap YAML viewer expanded with the five rules visible; (4) the EnforcementLog feed showing a DENY chip with prompt_injection_block and a ~2s ago timestamp. Caption strip: "60-second end-to-end run, all real artefacts on disk."